Establishing a Robust LLMOps Framework for Intelligent Automation: Strategies and Best Practices

1st Dasaprakash Krishnamurthy
Senior AI Architect
UST
London, UK
dasaprakash.krishnamurthy@ust.com

2nd Vinod Neelanath
Chief Product Officer
UST
Trivandrum, India
vinod.neelanath@ust.com

***Abstract*—This paper delves into the establishment of a robust LLMOps framework, crucial for intelligent automation in modern enterprises. LLMOps for intelligent automation platforms require composability and capability to use different architectures, embeddings, and Generative AI models seamlessly with domain focus. In this paper we discuss an architecture we designed for UST SmartOps, UST's Intelligent Automation platform, keeping automation at its core. Large Language Models (LLMs) have significantly advanced AI capabilities in natural language processing and understanding, but their deployment and management, integration and using them effectively for automation use cases pose unique challenges. LLMOps, similar to DevOps and MLOps, is dedicated to the operationalization of these models, addressing infrastructure, performance optimization, security, and continuous improvement. The paper outlines the strategies and best practices essential for LLMOps, emphasizing data management, model monitoring, and the integration of AI with automation to improve operational efficiency, reduce costs, and improve decision-making processes. Through systematic orchestration of tools, processes, and infrastructure, the document aims to provide a comprehensive guide to effectively leverage LLMs within intelligent automation frameworks.**

***Index Terms*—LLMOps, LLMOps framework, AI infrastructure, AI security, Model monitoring, LLM Cost Optimization**

I. Introduction

Large Language Models (LLMs) have revolutionized the field of Artificial Intelligence (AI) by providing unprecedented capabilities in natural language understanding, generation, and reasoning. However, deploying and managing these models at scale presents unique challenges related to infrastructure, performance optimization, security, monitoring, and continuous improvement. LLMOps (Large Language Model Operations) [1] [2], an emerging discipline similar to DevOps and MLOps [4] [6], focuses on operationalizing [12] the deployment, monitoring, and maintenance of LLMs to ensure that these powerful models can be used effectively in production environments. LLMOps is the systematic orchestration of processes, tools, and infrastructure to streamline the lifecycle of LLMs. In the rapidly evolving landscape of artificial intelligence and machine learning, establishing a robust LLMOps framework is integral to leveraging intelligent automation effectively.

979-8-3315-0754-1/25/$31.00 ©2025 IEEE

Intelligent automation holds paramount importance in modern enterprises due to its transformative potential across various operational dimensions. By integrating AI with automation, businesses can achieve higher efficiency, reduced operational costs, and improved accuracy in processes. This powerful combination enables enterprises to automate repetitive tasks, allowing human resources to focus on more strategic and creative endeavors [7]. In addition, intelligent automation improves decision making by providing real-time insights and data-driven analytics, thus fostering a more agile and responsive business environment.

LLMOps (Large Language Model Operations) [1] [2] generally focuses on managing, deploying, and optimizing large language models for various applications such as chatbots, content creation, and language translation. In contrast, LLMOps in intelligent automation specifically applies these models to automate business processes such as workflow automation, document processing, and robotic process automation (RPA). While LLMOps prioritizes scalability, accuracy, and performance across diverse use cases, LLMOps in intelligent automation emphasizes seamless integration with enterprise IT systems, ensuring models handle domain-specific data, compliance, and real-time decision making reliably.

The key differences also lie in objectives and metrics. LLMOps normally evaluates general performance through benchmarks like accuracy or latency, while LLMOps in intelligent automation focuses on business metrics such as process efficiency, ROI, and error reduction. It emphasizes task-specific fine-tuning, model interpretability, and close collaboration with business experts to meet specific automation goals. Ultimately, LLMOps in intelligent automation tailors LLMs to deliver measurable business value, making it distinct from the broader, more generalized approach of LLMOps. In intelligent automation, LLMOps emphasizes interoperability with RPA, low-code platforms and BPM toolsensuring that models not only generate insights but also trigger automated actions. These distinctions highlights how intelligent automation leverages LLMs to drive end-to-end process optimization while standard LLMOps focuses more on technical maintenance of the models

Large Language Models (LLMs) are increasingly becoming a core element of intelligent automation systems [8]

2025 Emerging Technologies for Intelligent Systems (ETIS) | 979-8-3315-0754-1/25/$31.00 ©2025 IEEE | DOI: 10.1109/ETIS64005.2025.10961869

Authorized licensed use limited to: UNIVERSITAETSBIBL STUTTGART. Downloaded on February 17,2026 at 03:44:35 UTC from IEEE Xplore. Restrictions apply.

[9]. However, managing LLMs within intelligent automation frameworks introduces unique challenges due to the scale, complexity, and dynamic nature of these models.

This paper outlines the strategies and best practices essential for developing and maintaining such a framework that can be applied to Intelligent Automation systems. This paper also explores the critical aspects of LLMOps, including architecture and best practices.

II. Key Challenges in implementing LLMOps

A. Data Management

Effectively managing data is a core aspect of successful LLMOps, but it brings numerous complexities that need addressing to ensure smooth operation and optimal performance of Large Language Models. Key challenges in data management for LLMOps include the collection of high-quality, diverse datasets, the preprocessing of this data to make it suitable for model training, and the efficient storage of massive datasets [1]. Furthermore, ensuring data security and compliance with regulations such as GDPR and CCPA, as well as implementing real-time data processing capabilities, are critical components for maintaining the efficacy and relevance of LLMs in intelligent automation frameworks [9].

B. Model Monitoring

Deploying and monitoring Large Language Models (LLMs) in LLMOps introduces several intricate challenges. The sheer size and complexity of LLMs require robust infrastructure capable of handling substantial computational loads and ensuring high availability. Performance optimization is crucial to meet the latency and throughput demands of real-time applications. Continuous monitoring is essential to detect and mitigate drifts in model performance, automation percentage and ensuring that the models remain accurate and reliable over time. Moreover, implementing feedback loops for ongoing improvement and scalability presents another layer of complexity, necessitating advanced tools and methodologies to seamlessly integrate updates without disrupting the operational workflow.

C. Scale

LLMs can range from hundreds of millions to trillions of parameters, requiring massive computational resources. These models demand advanced hardware, such as GPUs and TPUs, to handle the immense computational load and ensure efficient training and inference. The infrastructure must support the high memory bandwidth and storage requirements of these models, often necessitating distributed computing environments. Additionally, the energy consumption associated with training and deploying LLMs is substantial, posing environmental and operational cost concerns that must be addressed. Therefore, optimizing the scale and efficiency of LLM operations is paramount to sustainable and effective AI deployment for real-time or near real-time inferences to keep up with business SLAs [5].

D. Security and compliance

LLMs pose data security and privacy risks including unintentional leakage of sensitive information or susceptibility to adversarial attacks. A significant challenge is the phenomenon of hallucinations, where LLMs generate plausible but incorrect information. Implementing guardrails, such as strict validation checks and human-in-the-loop systems, can mitigate the risks associated with these hallucinations and ensure the reliability of the information produced by LLMs. These safeguards are vital for maintaining the quality and accuracy of the output, particularly in sensitive applications.

III. Strategies for building a robust LLMOps for Intelligent Automation

The challenges described in the previous section require robust strategies including modular architectures for seamless model accessibility and updates, continuous monitoring for data and model drift, and automated data pipelines to enhance data governance. This paper explains the architecture we used for building LLMOps catering to Intelligent Automation and the strategies we followed to mitigate the challenges.

A. Enabling Scalable Deployment and Integration

LLMOps facilitates the seamless integration of LLMs into existing automation platforms and workflows. By providing robust infrastructure for scalable deployment, LLMOps ensures that LLMs can be effectively utilized across diverse use cases and at different levels of demand, ranging from small-scale tasks to enterprise-level applications. Figure 1 demonstrates the architecture and the various components required for LLMOps in an Intelligent Automation platform [10]. This architecture provides interface to query LLMs directly and through RAG and agents.

- **LLMOps pipeline:** The architecture for LLMOps pipeline presented here details how we integrated and utilized different types of generative models (cloud based and self-hosted Large Language Models, Large Vision Models and Large Multi-modal Models) in Intelligent Automation use cases. The pipeline automates model training, fine-tuning, adapting, deployment, monitoring and real-time serving. The multi-agent framework facilitates complex interactions and integrates different components required for agentic orchestration. Caching mechanisms, such as response and token-level caching, reduce latency by storing frequently used outputs and optimizing computation. To maintain control over inputs and generated outputs, guardrails like content filtering, denied topics, word and sensitive information filters are included. LLMOps tools for configuration management, monitoring, dashboards, cost control, quota management, response time tracking, evaluation and auditing are built for supporting these models. LLM Gateway provides integration with RPA tools, Intelligent Document Processing Tools and external applications and takes care of data ingestion, connectivity, security and serialization

Authorized licensed use limited to: UNIVERSITAETSBIBL STUTTGART. Downloaded on February 17,2026 at 03:44:35 UTC from IEEE Xplore. Restrictions apply.

which are exclusively designed for supporting Intelligent Automation use cases.

- **Multi-cloud and Hybrid Deployment:** LLMOps allows for flexible deployment models, supporting multi-cloud environments and hybrid setups (on-premise + cloud). This is essential for scaling LLM-based automation solutions across different business functions or geographical regions [11].
- **API Integration and Orchestration:** LLMOps ensures that LLMs can be integrated with automation systems via APIs or microservices, enabling seamless interactions between the language models and other automation tools such as robotic process automation (RPA), workflow engines, and decision management systems.

B. Continuous Model Optimization and Adaptation

In intelligent automation, models must continuously evolve to adapt to changing business requirements, regulatory environments, or user preferences. LLMOps provides the necessary tools and processes to enable this continuous optimization and fine-tuning of LLMs in production environments.

- **Model Fine-tuning:** LLMOps allows for domain-specific fine-tuning, where models can be tailored to industry-specific needs, such as healthcare, finance, or legal automation. For instance, LLMs can be fine-tuned to generate domain-relevant content or provide industry-specific insights.
- **Active Learning and Feedback Loops:** LLMOps supports integrating human-in-the-loop (HITL) mechanisms, where feedback from end-users or experts can be incorporated into the model's fine-tuning process. This ensures continuous improvement of model performance and relevance, which is critical for adaptive intelligent automation systems.

C. Monitoring and Error Handling in Automated Processes

In intelligent automation, where processes must be reliable and efficient, monitoring and error detection are crucial. LLMOps ensures that LLM-driven systems maintain optimal performance through real-time monitoring, error detection, and performance optimization.

- **Anomaly Detection and Alerting:** By monitoring LLM performance metrics such as latency, accuracy, and throughput in addition to business SLAs, LLMOps helps detect anomalies or deviations from expected behavior. This is vital for ensuring that LLMs do not introduce errors or delays into automated workflows.
- **Error Handling and Correction:** Intelligent automation systems often involve complex decision-making tasks. LLMOps ensures that when an LLM makes a mistake, appropriate error-handling mechanisms are in place to either correct the output automatically or escalate it to human operators for intervention.

D. Ensuring Security, Compliance, and Governance

As LLMs become a central part of intelligent automation systems, security, compliance, and governance become increasingly important. LLMOps provides the frameworks necessary to ensure that LLMs operate within the boundaries of security protocols and regulatory standards.

- **Data Privacy and Compliance:** LLMOps ensures that LLMs comply with privacy regulations like GDPR or CCPA by controlling how sensitive data is handled within automated processes. Guard railing techniques such as data anonymization, encryption, and secure model deployment are employed to protect user data.
- **Auditability and Traceability:** LLMOps enables the tracking of model outputs, updates, and interactions, ensuring that all decisions made by the model within an intelligent automation system are audited. This is critical for industries such as finance or healthcare, where regulatory oversight requires detailed accountability of AI-driven decisions.

E. Cost Efficiency and Resource Management

LLMOps is essential for optimizing the cost and resource allocation when using LLMs in intelligent automation, given the high computational requirements of these models. Efficient resource management ensures that the use of LLMs in automation systems remains cost-effective and scalable.

- **Auto-scaling and Load Management:** LLMOps enables the automatic scaling of compute resources or quota management based on demand. For example, during peak loads, more computational resources can be allocated, and during off-peak times, the system can scale down to minimize costs.
- **Optimization Techniques:** LLMOps incorporates techniques such as model pruning (removing redundant parameters) and quantization (reducing the precision of model parameters) to reduce the computational and memory footprint of LLMs. This ensures that the models can perform efficiently within the resource constraints of an automated environment.

F. Enhancing Responsiveness and User Experience

Intelligent automation systems are often user-facing, and LLMs are frequently deployed in conversational agents, virtual assistants, and decision support systems. LLMOps ensures that these models respond in real time and provide high-quality, contextually relevant outputs, enhancing the overall user experience.

- **Low-latency Response Management:** LLMOps employs caching, parallelization, and optimized hardware configurations (e.g., using GPUs or TPUs) to ensure that LLMs deliver responses in real-time, even when handling complex queries.
- **Custom User Interaction Models:** LLMOps enables the customization of LLMs for specific user interaction scenarios. For example, LLMs can be fine-tuned to handle

Authorized licensed use limited to: UNIVERSITAETSBIBL STUTTGART. Downloaded on February 17,2026 at 03:44:35 UTC from IEEE Xplore. Restrictions apply.

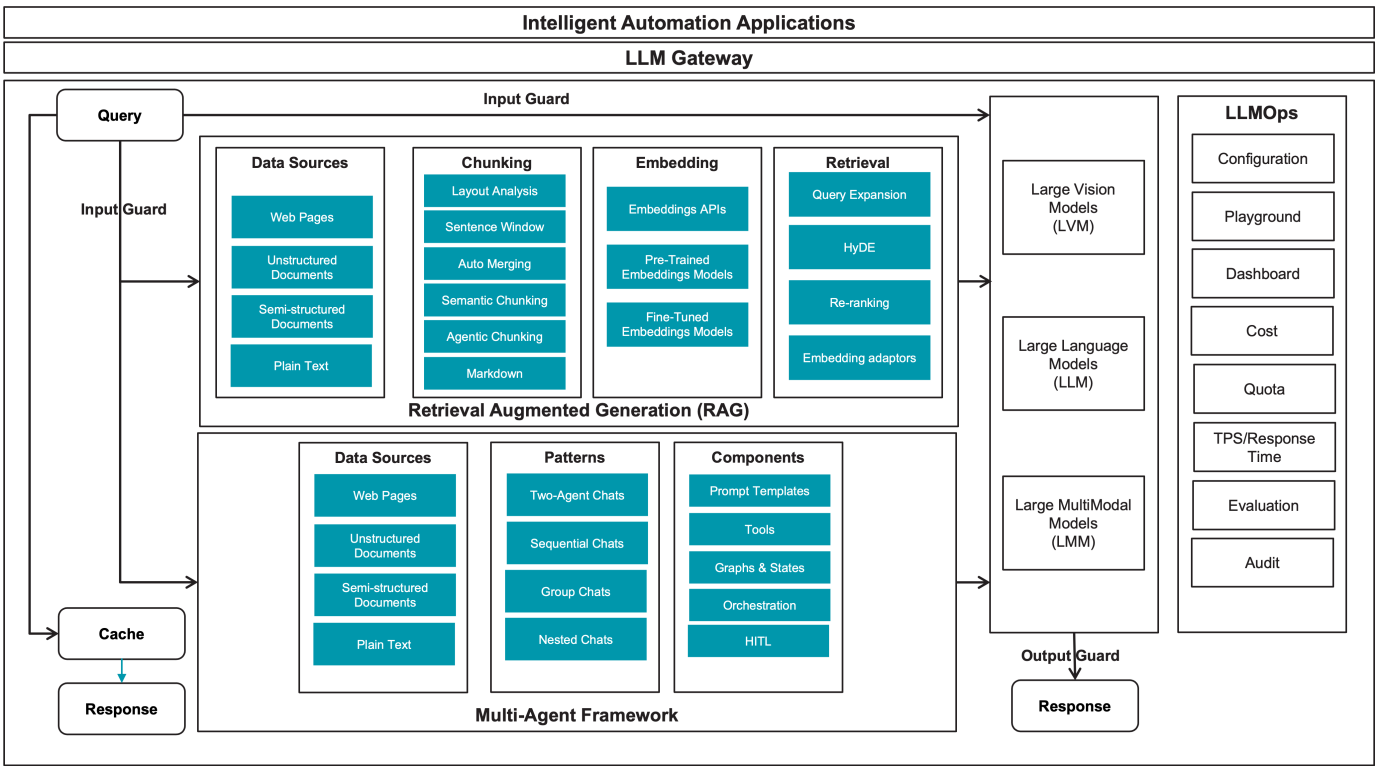

Fig. 1. LLMOps Architecture

customer service queries in a way that aligns with a company's organizational policies, improving both user satisfaction and operational efficiency.

G. Supporting Multi-modal and Multi-agent automations

For automation in complex automation systems LLMs may not work in isolation but in conjunction with other AI models (e.g., computer vision, decision trees) or automation agents. LLMOps plays a vital role in orchestrating these multi-modal systems and ensuring that LLMs integrate smoothly with other intelligent agents.

- **Cross-modal Integration:** LLMOps enables the combination of LLMs with other AI modalities, such as vision or speech recognition, to create more comprehensive intelligent automation systems. For instance, an LLM could generate a response based on an image that has been processed by a computer vision model.
- **Multi-agent Systems:** In some intelligent automation scenarios, multiple AI agents must work together to complete tasks. LLMOps ensures that these agents communicate effectively, share context, and work with synergy to achieve common objectives.

IV. LLMOps in the Context of Retrieval-Augmented Generation (RAG) and Agentic AI

As Large Language Models (LLMs) continue to evolve, new approaches like Retrieval-Augmented Generation (RAG) and Agentic AI are emerging to push the boundaries of their capabilities. These techniques introduce unique challenges and opportunities for LLMOps, which focuses on operationalizing LLMs at scale.

A. Retrieval-Augmented Generation (RAG)

Enhancing LLMs with External Knowledge RAG combines the power of LLMs with external knowledge bases, making it possible for models to generate more accurate, contextually relevant, and up-to-date responses. Instead of relying solely on the information encoded during pre-training, RAG enables LLMs to retrieve real-time data or specific facts from external sources (e.g., databases, search engines) during the generation process.

From an LLMOps perspective, RAG introduces additional layers of complexity in managing LLM workflows:

- **Dynamic Data Retrieval:** RAG systems require efficient integration between LLMs and external databases or knowledge sources. LLMOps must ensure that these integrations are reliable and scalable, supporting real-time queries without introducing latency.
- **Knowledge Freshness and Accuracy:** One of the main benefits of RAG is that it allows LLMs to generate responses based on current data, overcoming the limitations of static pre-trained models. LLMOps teams must implement mechanisms to continuously update and curate external knowledge sources to ensure the LLMs are drawing from accurate and up-to-date information.
- **Model Monitoring:** Because RAG systems rely on both the LLM and the external retrieval system, monitoring

Authorized licensed use limited to: UNIVERSITAETSBIBL STUTTGART. Downloaded on February 17,2026 at 03:44:35 UTC from IEEE Xplore. Restrictions apply.

becomes more complex. LLMOps must track not only the model's performance but also the quality of the retrieved information. For instance, in customer support automation, a RAG system can query product manuals or knowledge bases in real time to provide precise answers to customer inquiries. LLMOps teams need to monitor both retrieval accuracy and model behavior to ensure high-quality responses across diverse customer needs.

B. Agentic AI

Empowering LLMs with Autonomous Decision-Making Agentic AI refers to LLM-powered agents that can autonomously interact with environments, make decisions, and execute tasks based on user instructions or goals. These agents combine LLMs with other AI systems and tools (e.g., APIs, RPA tools, search engines) to perform complex tasks autonomously, effectively acting as intelligent agents that go beyond just generating text.

LLMOps plays a pivotal role in the operational success of Agentic AI by addressing several key challenges:

- **Autonomous Task Execution:** Agentic AI systems involve not only natural language understanding but also the ability to take action. LLMOps must manage how these agents interact with external systems (like APIs or robotic process automation tools) in a reliable, secure, and efficient manner.
- **Real-time Monitoring and Safety:** Given that Agentic AI systems are capable of autonomous decision-making, real-time monitoring is crucial to ensure that the agents operate within predefined safety and ethical boundaries. LLMOps must establish guardrails to prevent agents from making harmful or erroneous decisions, and alert human operators when necessary.
- **Human-in-the-Loop (HITL):** For complex tasks or situations where an agent's autonomy needs to be limited, LLMOps can implement human-in-the-loop processes, allowing humans to intervene or guide the agent's decisions. This is especially useful in high-stakes environments, such as healthcare or legal services, where autonomous decisions can have significant consequences. In an Agentic AI scenario, an LLM-powered agent could autonomously draft a contract, validate details via external APIs (like legal databases), and suggest clauses to users. LLMOps ensures that the agent acts within defined parameters, monitors its output for errors, and escalates to human reviewers when necessary.

Both RAG and Agentic AI expand the functional scope of LLMs, pushing them beyond static text generation into real-time, action-oriented domains. LLMOps ensures that these advanced systems are deployed effectively, with robust monitoring, security, and scalability. It also provides a framework to continuously improve the performance of LLMs as they interact with dynamic environments, external knowledge bases, and autonomous task execution systems.

RAG enhances Agentic AI by providing timely, relevant information for decision-making, while Agentic AI leverages that information to autonomously execute tasks. Together, they enable intelligent automation to be both knowledge-driven and action-oriented.

In an automated insurance claims process, RAG retrieves relevant policy documents and previous claim records, while Agentic AI evaluates the claim, communicates with the claimant and processes the payout decision.

V. Conclusion

The pivotal role of LLMOps in Intelligent Automation is critical to making large language models a reliable and scalable component of intelligent automation systems. By providing the infrastructure, tooling, and processes necessary to deploy, monitor, and continuously improve LLMs, LLMOps ensures that these models can be operationalized effectively in real-world automation scenarios. As organizations increasingly rely on intelligent automation to optimize their workflows, LLMOps will continue to play a central role in ensuring that LLMs deliver high performance, security, and cost-efficiency at scale.

References

[1] Diaz de Arcaya, Josu and López-De-Armentia, Juan and Miñón, Raúl and Ojanguren, Iker and Torre-Bastida, Ana-Isabel, "Large Language Model Operations (LLMOps): Definition, Challenges, and Lifecycle Management," 2024 9th International Conference on Smart and Sustainable Technologies (SpliTech), Bol and Split, Croatia, 2024, pp. 1-4, doi: 10.23919/SpliTech61897.2024.10612341.

[2] R. Shan and T. Shan, "Enterprise LLMOps: Advancing Large Language Models Operations Practice," 2024 IEEE Cloud Summit, Washington, DC, USA, 2024, pp. 143-148, doi: 10.1109/CloudSummit61220.2024.00030.

[3] T. Chen, "Challenges and Opportunities in Integrating LLMs into Continuous Integration/Continuous Deployment (CI/CD) Pipelines," 2024 5th International Seminar on Artificial Intelligence, Networking and Information Technology (AINIT), Nanjing, China, 2024, pp. 364-367, doi: 10.1109/AINIT61980.2024.10581784.

[4] Dominik Kreuzberger and Niklas Kühl and Sebastian Hirschl, "Machine Learning Operations (MLOps): Overview, Definition, and Architecture," 2022 https://arxiv.org/abs/2205.02302.

[5] Lucas Baier, Fabian Jöhren, and Stefan Seebacher, "Challenges in the deployment and operation of machine learning in practice,"

[6] Singla, Amandeep. (2023). "Machine Learning Operations (MLOps): Challenges and Strategies. Journal of Knowledge Learning and Science Technology" ISSN: 2959-6386 (online). 2. 333-340. 10.60087/jklst.vol2.n3.p340.

[7] Kulkarni, Akshay, Shivananda, Adarsha, Kulkarni, Anoosh, Gudivada, Dilip. (2023). "LLMs for Enterprise and LLMOps."

[8] S. Stoykova, R. Hrischev and N. Shakev, "Intelligent Robotic Process Automation for Small and Medium-sized Enterprises," 2022 International Conference Automatics and Informatics (ICAI), Varna, Bulgaria, 2022, pp. 223-228, doi: 10.1109/ICAI55857.2022.9960077.

[9] O. C. Williams and F. Olajide, "Towards the Design of an Intelligent Automation Framework for Business Processes," 2022 5th International Conference on Information and Computer Technologies (ICICT), New York, NY, USA, 2022, pp. 13-17, doi: 10.1109/ICICT55905.2022.00010.

[10] I. Arawjo, C. Swoopes, P. Vaithilingam, M. Wattenberg and E. Glassman, "Chainforge: A visual toolkit for prompt engineering and llm hypothesis testing", arXiv preprint, 2023.

[11] Paul Singh; Anurag Karuparti; John Maeda, "Generative AI for Cloud Solutions: Architect modern AI LLMs in secure, scalable, and ethical cloud environments" , Packt Publishing, 2024.

[12] A. Bodor, M. Hnida and D. Najima, "From Development to Deployment: An Approach to MLOps Monitoring for Machine Learning Model Operationalization," 2023 14th International Conference on Intelligent Systems: Theories and Applications (SITA), Casablanca, Morocco, 2023, pp. 1-7, doi: 10.1109/SITA60746.2023.10373733.

Authorized licensed use limited to: UNIVERSITAETSBIBL STUTTGART. Downloaded on February 17,2026 at 03:44:35 UTC from IEEE Xplore. Restrictions apply.